

# E0 - 243 Computer Architecture Assignment - 2

Mandar Bhalerao - mandarmanoj@iisc.ac.in

Kunal Pansare - kunaldinesh@iisc.ac.in

November 26, 2023

## 1 Introduction

We are required to optimize the problem of "Dilated Convolution(DC)" which involves an input matrix and a kernel matrix. The output of the dilated convolution is stored in the output matrix. The dimensions of these matrix are:-

### Input:

1. Input Matrix: Input\_Row x Input\_Column
2. Kernel Matrix: Kernel\_Row x Kernel\_Column

### Output:

Output Matrix:  $(\text{Input\_Row} - \text{Kernel\_Row} + 1) \times (\text{Input\_Column} - \text{Kernel\_Column} + 1)$

We will Optimize the Single Threaded Dilated Convolution and Implement and Optimize Multi-Threaded Dilated Convolution of the given unoptimized implementation provided on GitHub.

## 2 Optimization Methods used and Performance Analysis

### 2.1 Optimization Methods

**I. Loop Unrolling:-** Loop unrolling is a compiler optimization technique that aims to improve the performance of a program by reducing the overhead of loop control code. The idea is to process multiple iterations of a loop in a single iteration, effectively reducing the number of loop-control instructions and improving instruction-level parallelism.

**II. Elimination of Redundant Computations:-** Redundant computations refer to performing the same calculation multiple times when the result remains unchanged. Some strategies used are:-

- a) Storing Intermediate Results
- b) Loop Invariant Code Motion
- c) Common Subexpression Elimination
- d) Use of Temporary Variables

**III. Strength Reduction:-** Strength reduction is a compiler optimization technique that involves replacing expensive operations with equivalent, but less expensive, operations. The goal is to improve the overall performance of a program by reducing the computational cost of certain expressions.

**IV. Single Instruction, Multiple Data(SIMD):-** SIMD stands for Single Instruction, Multiple Data. It's a type of parallel computing architecture that performs the same operation on multiple data elements simultaneously. In simpler terms, SIMD allows a processor to apply a single instruction to multiple pieces of data at the same time, improving the overall efficiency of computations.

### 2.2 Performance Analysis

Given below is the Performance Analysis of obtaining the output matrix when it is executed for combinations of input matrix dimensions and kernel matrix dimensions respectively. The Performance Metrics used for analysis are:-

## [PartA-I] Optimize single-threaded DC(CPU)

**a. Improvement in Single Threded Execution(%)** :- The metric calculates the percentage reduction in execution time achieved by the single-threaded implementation compared to a reference execution. This metric quantifies the efficiency gain achieved by the single-threaded approach, expressing it as a percentage improvement relative to the reference execution time.

$$\text{Improvement}(\%) = (1 - (\text{Single Thread Execution Time} / \text{Reference Execution Time})) * 100$$

**b. Speedup in Single Threded Execution** :- Speedup is a performance metric that quantifies the improvement in the execution time of a program or algorithm when compared to a reference or baseline execution. It is typically expressed as the ratio of the reference execution time to the improved execution time. The speedup is calculated using the formula:

$$\text{Speedup} = \text{Reference Execution Time} / \text{Single Thread Execution Time}$$

| Sr. No | Kernel Matrix dimensions | Reference Execution Time(ms) | Single Thread Execution Time(ms) | Imp in STE(%) | Speedup(STE) |
|--------|--------------------------|------------------------------|----------------------------------|---------------|--------------|
| 1      | 3x3                      | 925.2                        | 364.7                            | 60.57         | 2.53         |
| 2      | 5x5                      | 2427.8                       | 443.1                            | 81.74         | 5.47         |
| 3      | 7x7                      | 4705.6                       | 795.7                            | 83.08         | 5.91         |
| 4      | 9x9                      | 7779.0                       | 1253.3                           | 83.88         | 6.20         |
| 5      | 11x11                    | 11586.3                      | 1848.6                           | 84.0          | 6.26         |
| 6      | 13x13                    | 16511.1                      | 2517.71                          | 84.75         | 6.55         |
| 7      | 64x64                    | 352101.0                     | 63791.7                          | 81.88         | 5.51         |

Table 1: Input Matrix with dimensions 4096x4096




Figure 1

| Sr. No | Kernel Matrix dimensions | Reference Execution Time(ms) | Single Thread Execution Time(ms) | Imp in STE(%) | Speedup(STE) |
|--------|--------------------------|------------------------------|----------------------------------|---------------|--------------|
| 1      | 3x3                      | 3711.4                       | 1489.8                           | 59.85         | 2.49         |
| 2      | 5x5                      | 10019                        | 1767.7                           | 82.35         | 5.66         |
| 3      | 7x7                      | 19569.1                      | 3275.9                           | 83.25         | 5.97         |
| 4      | 9x9                      | 31386.6                      | 5043.3                           | 83.93         | 6.22         |
| 5      | 11x11                    | 46482.1                      | 7350                             | 84.18         | 6.32         |
| 6      | 13x13                    | 66888.1                      | 10088.4                          | 84.91         | 6.63         |
| 7      | 64x64                    | 1431110                      | 247423                           | 82.71         | 5.78         |

Table 2: Input Matrix with dimensions 8192x8192




Figure 2

| Sr. No | Kernel Matrix dimensions | Reference Execution Time(ms) | Single Thread Execution Time(ms) | Imp in STE(%) | Speedup(STE) |
|--------|--------------------------|------------------------------|----------------------------------|---------------|--------------|
| 1      | 3x3                      | 14921.3                      | 5886.6                           | 60.54         | 2.53         |
| 2      | 5x5                      | 40086.5                      | 7028.3                           | 82.46         | 5.70         |
| 3      | 7x7                      | 77150.9                      | 12828.1                          | 83.37         | 6.01         |
| 4      | 9x9                      | 125654                       | 19902.6                          | 84.16         | 6.31         |
| 5      | 11x11                    | 186615                       | 28219                            | 84.87         | 6.61         |
| 6      | 13x13                    | 267618                       | 38078                            | 85.77         | 7.02         |

Table 3: Input Matrix with dimensions 16384x16384

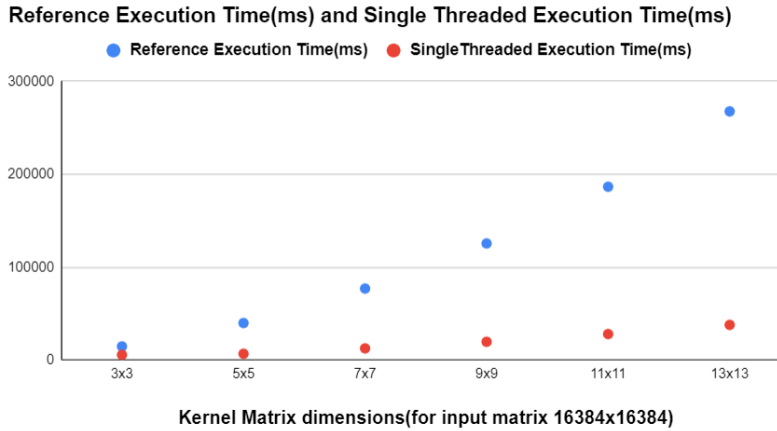


Figure 3

## 2.3 Analysis using perf tools




Figure 4




Figure 5




Figure 6

## [PartA-II] Optimize multi-threaded DC(CPU)

It refers to implementing dilated convolution in a multi-threaded programming paradigm. Multi-threading involves dividing a program into smaller threads that can run concurrently, potentially improving performance on multi-core processors.

| Sr. No | Kernel Matrix dimensions | Reference Execution Time(ms) | Multi Thread Execution Time(ms) | Imp in MTE(%) | Speedup(MTE) |
|--------|--------------------------|------------------------------|---------------------------------|---------------|--------------|
| 1      | 3x3                      | 925.2                        | 364.7                           | 60.57         | 2.53         |
| 2      | 5x5                      | 2427.8                       | 443.1                           | 81.74         | 5.47         |
| 3      | 7x7                      | 4705.6                       | 795.7                           | 83.08         | 5.91         |
| 4      | 9x9                      | 7779.0                       | 1253.3                          | 83.88         | 6.20         |
| 5      | 11x11                    | 11586.3                      | 1848.6                          | 84.0          | 6.26         |
| 6      | 13x13                    | 16511.1                      | 2517.71                         | 84.75         | 6.55         |
| 7      | 64x64                    | 352101.0                     | 63791.7                         | 81.88         | 5.51         |

Table 4: Input Matrix with dimensions 4096x4096

| Number of Threads<br>Input Matrix Dim. | Reference(ms) | 4       | 8      | 12     | 16     |
|--|---------------|---------|--------|--------|--------|
| 1024                                   | 998.1         | 48.27   | 46.49  | 39.85  | 47.65  |
| 2048                                   | 4039.8        | 356.16  | 339.31 | 347.46 | 281.78 |
| 4096                                   | 16511         | 1679.7  | 1754.2 | 1850   | 1830   |
| 8192                                   | 66888.1       | 6590.42 | 6906.5 | 7530   | 7490.9 |

Table 5: Kernel Matrix with dimensions 13x13




Figure 7

| Number of Threads<br>Input Matrix Dim. | Reference(ms) | 4       | 8       | 12     | 16      |
|--|---------------|---------|---------|--------|---------|
| 1024                                   | 20996         | 1595.4  | 1445.3  | 1312.2 | 1379.7  |
| 2048                                   | 90530.2       | 8682.64 | 7066.9  | 7033   | 6159.3  |
| 4096                                   | 352101        | 40590.5 | 40415.4 | 44033  | 43298.3 |
| 8192                                   | 1431110       | 157475  | 165831  | 176217 | 175417  |

Table 6: Kernel Matrix with dimensions 64x64




Figure 8

### 3 Conclusion

#### 3.1 Conclusion for Optimized Single-Threaded Dilated Convolution

In conclusion, the optimized single-threaded dilated convolution implementation has demonstrated notable improvements in performance. By strategically incorporating loop unrolling, memory access optimizations, Elimination of redundant computations, SIMD, Strength Reduction we have significantly reduced the execution time. This optimized version maintains the accuracy of the convolution operation while leveraging computational resources more efficiently on a single processing core.

#### 3.2 Conclusion for Multi-Threaded Dilated Convolution

The multi-threaded dilated convolution implementation has demonstrated the power of parallelism in accelerating the computation of dilated convolutions. By dividing the workload across multiple threads, we harnessed the capabilities of modern multi-core processors to achieve substantial reductions in execution time.